

Audio key estimation of digital music with CNNs

Udacity Machine Learning Nanodegree - Capstone project

Table of Contents

The project is structured as stated in section 'Project Design' of the Capstone project proposal.

—

Data Preprocessing

Million Song Dataset - selection of appropriate songs, separate jupyter notebook

Signal Processing and Feature Extraction - separate jupyter notebook

Model Preparation

Load and prepare data - read spectrogram images, conversion to tensors

Splitting data into training/testing sets

CNN model architecture

CNN model parameter

Model Training and Evaluation

Model training

Model evaluation and comparison

Current version of the project is working, but

the project is still ongoing...

discussion and remarks of what to do can be found in section

reasons / todo

Data Preprocessing

Million Song Dataset

- utilized to select appropriate song samples
- holds information about key and mode per song (targets)

Jupyter Notebook [msd \(./00.hlp/msd/msd.ipynb\)](#)

outputs: csv file *songs_conf=75_tracks_filt.csv*, which holds all songs with key confidence and mode confidence > 0.75

```
In [1]: # LIST SELECTED SONGS
import os
import pandas as pd
from IPython.display import display

selsongsfile = os.path.join ('00.hlp', 'msd', 'songs_conf=75_tracks_filt.csv')
selsongs = pd.read_csv (selsongsfile, header=0, index_col=0)
display (selsongs.head (1))
print ('[i] number of records:', len (selsongs))
```

| | key | key_confidence | mode | mode_confidence | track_id | song_id | artist_name | song_title |
|---|-----|----------------|------|-----------------|--------------------|--------------------|------------------|---------------|
| 0 | 7 | 0.896 | 1 | 0.852 | TRMMMGL128F92FD6AB | SOHSSPG12A8C144BE0 | Clifford T. Ward | Mad About You |

[i] number of records: 47913

```
In [2]: # LOAD AUDIO DATASET
import os
import numpy as np
from sklearn import datasets

PARAM_RND_STATE = 42

container_path = os.path.join ('src_audio')
load_content = False
description = ['key C, mode minor', 'key C, mode major',
               'key C#, mode minor', 'key C#, mode major',
               'key D, mode minor', 'key D, mode major',
               'key D#, mode minor', 'key D#, mode major',
               'key E, mode minor', 'key E, mode major',
               'key F, mode minor', 'key F, mode major',
               'key F#, mode minor', 'key F#, mode major',
               'key G, mode minor', 'key G, mode major',
               'key G#, mode minor', 'key G#, mode major',
               'key A, mode minor', 'key A, mode major',
               'key A#, mode minor', 'key A#, mode major',
               'key B, mode minor', 'key B, mode major']

src_audio_data = datasets.load_files (container_path=container_path,
                                     description=description,
                                     load_content=load_content,
                                     random_state=PARAM_RND_STATE)
```

```
In [3]: # FYI: LIST SOME OF THE USED SONGS
filenames = list (os.path.basename (filepath) for filepath in src_audio_data['filenames'])
usedsongs_track_id = list (os.path.splitext (fn)[0] for fn in filenames)
usedsongs = selsongs.query ('track_id in @usedsongs_track_id')

display (usedsongs.sample(5))
print ('[i] min of: key_confidence =', usedsongs['key_confidence'].min (), ', ', '\
      'mode_confidence =', usedsongs['mode_confidence'].min ())
print ('[i] number of records:', len (usedsongs))
```

| | key | key_confidence | mode | mode_confidence | track_id | song_id | artist_name | song_title |
|-------|-----|----------------|------|-----------------|--------------------|--------------------|---------------------|-----------------|
| 10070 | 0 | 1.0 | 0 | 0.849 | TRRTCMy128F9326B32 | SOYGPIR12AB018646A | The Levon Helm Band | A Fool In Love |
| 22522 | 1 | 1.0 | 1 | 0.927 | TRNFTSL128F4259524 | SOAAUTE12A8C13491F | Vanilla Sky | Gotta Believe |
| 25187 | 4 | 1.0 | 1 | 1.000 | TRPJEEr128F426E291 | SOLUZSQ12A8C13C968 | Embrace | Anywhere You Go |
| 46978 | 3 | 1.0 | 0 | 0.810 | TRYNAXP128F93205A9 | SOGZYYE12AB0183215 | CJ Stone | Shining Star |
| 45681 | 6 | 1.0 | 0 | 0.913 | TRKVLFA12903CD12C4 | SOQZPLI12AB0186D42 | temposhark | Bye Bye Baby |

[i] min of: key_confidence = 0.809 , mode_confidence = 0.777
[i] number of records: 240

Signal Processing and Feature Extraction

- create spectrograms of audio files with discrete Fourier transform (DFT)
- save spectrograms as images for further use in CNN

Jupyter Notebook [fft \(/00.hlp/fft.ipynb\)](#)

outputs: spectrograms (png images) of audio files with same folder structure as *src_audio* in new container path named *src_spectro*

Example of a spectrogram image

Model Preparation

Load and prepare data

```
In [4]: # LOAD SPECTROGRAM FILENAMES
import os
import numpy as np
from sklearn import datasets

PARAM_RND_STATE = 42

container_path = os.path.join ('src_spectro')
load_content = False
description = ['key C, mode minor', 'key C, mode major',
              'key C#, mode minor', 'key C#, mode major',
              'key D, mode minor', 'key D, mode major',
              'key D#, mode minor', 'key D#, mode major',
              'key E, mode minor', 'key E, mode major',
              'key F, mode minor', 'key F, mode major',
              'key F#, mode minor', 'key F#, mode major',
              'key G, mode minor', 'key G, mode major',
              'key G#, mode minor', 'key G#, mode major',
              'key A, mode minor', 'key A, mode major',
              'key A#, mode minor', 'key A#, mode major',
              'key B, mode minor', 'key B, mode major']

src_spectro_data = datasets.load_files (container_path=container_path,
                                       description=description,
                                       load_content=load_content,
                                       random_state=PARAM_RND_STATE)

src_spectro_data.keys ()
```

Out[4]: dict_keys(['filenames', 'target_names', 'DESCR', 'target'])

```
In [5]: print ('[i] example of loaded spectrogram file data:')
print ('    spectrogram image name:', src_spectro_data['filenames'][0])
print ('    spectrogram image key-mode pair:',\
      src_spectro_data['target_names'][src_spectro_data['target'][0]],\
      '=', src_spectro_data['DESCR'][src_spectro_data['target'][0]],\
      '= target class', src_spectro_data['target'][0])

[i] example of loaded spectrogram file data:
spectrogram image name: src_spectro/1-0/TRLZZOJ128F1494C12.png
spectrogram image key-mode pair: 1-0 = key C#, mode minor = target class 2
```

Read in images, convert to tensors

Keras Conv2D layers expect a **4D tensor with shape (batch, rows, cols, channels)** (if param `data_format='channels_last'`) (src: [Keras Conv2D](https://keras.io/layers/convolutional/#conv2d) (<https://keras.io/layers/convolutional/#conv2d>))

```
In [6]: # open a random image and take a look at the attributes
import numpy as np
from PIL import Image

im = Image.open (src_spectro_data['filenames'][0])
print ('[i] image size:', im.size)
print ('[i] pixel format:', im.mode)

[i] image size: (162, 162)
[i] pixel format: RGB
```

images are of size (162, 162) and have 3 channels

for CNN: no need to change target size

below functions read images and convert those to tensors - code taken from Udacity MLND dog-project

```
In [7]: from keras.preprocessing import image
from tqdm import tqdm

def path_to_tensor (img_path):
    # loads RGB image as PIL.Image.Image type
    img = image.load_img (img_path, color_mode='grayscale')#, target_size=(88, 88))
    # convert PIL.Image.Image type to 3D tensor
    x = image.img_to_array (img)
    # convert 3D tensor to 4D tensor
    return np.expand_dims (x, axis=0)

def paths_to_tensor (img_paths):
    list_of_tensors = [path_to_tensor (img_path) for img_path in tqdm (img_paths)]
    return np.vstack (list_of_tensors)
```

Using TensorFlow backend.

```
In [8]: from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

spectro_tensors = paths_to_tensor (src_spectro_data['filenames'])#.astype ('float32') / 255

100%|██████████| 240/240 [00:00<00:00, 1665.64it/s]
```

```
In [9]: print ('[i] shape of spectrogram tensors:', spectro_tensors.shape)

[i] shape of spectrogram tensors: (240, 162, 162, 1)
```

```
In [10]: from keras.utils import np_utils
targets = np_utils.to_categorical (np.array (src_spectro_data['target']), 24)
print ('[i] number of output classes:', targets.shape[1])

[i] number of output classes: 24
```

Split data into train and test set

```
In [11]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = \
    train_test_split (spectro_tensors, targets, test_size=0.10, shuffle=True, random_state=PARAM_RND_STATE)

print ('[i] Training dataset consists of {} samples'.format (X_train.shape[0]))
print ('[i] Testing dataset consists of {} samples'.format (X_test.shape[0]))

[i] Training dataset consists of 216 samples
[i] Testing dataset consists of 24 samples
```

Model architecture

```
In [12]: from keras import layers, models
from keras import backend as K

# clear everything known of past instances ("useful to avoid clutter from old models / layers")
K.clear_session ()

# input layer
inputs = layers.Input (shape=spectro_tensors.shape[1:], name='input')

# hidden layers
net = layers.Conv2D (filters=8, kernel_size=(1,1), strides=(1,1),
                    padding='same', # don't lose information due to conv window runs out of image / stride
                    s = 1 = OK
                    activation='relu',
                    name='conv2d_1') (inputs)
#net = layers.MaxPooling2D (pool_size=(2,2), strides=None, name='maxp_1') (net)

net = layers.Conv2D (filters=16, kernel_size=(2,2), strides=(1,1),
                    padding='same',
                    activation='relu',
                    name='conv2d_2') (net)
net = layers.MaxPooling2D (pool_size=(2,2), strides=None, name='maxp_2') (net)

net = layers.Conv2D (filters=32, kernel_size=(2,2), strides=(1,1),
                    padding='same',
                    activation='relu',
                    name='conv2d_3') (net)
net = layers.MaxPooling2D (pool_size=(2,2), strides=None, name='maxp_3') (net)

net = layers.Conv2D (filters=64, kernel_size=(2,2), strides=(1,1),
                    padding='same',
                    activation='relu',
                    name='conv2d_4') (net)
net = layers.MaxPooling2D (pool_size=(2,2), strides=None, name='maxp_4') (net)

# 'flatten layer'
net = layers.GlobalAveragePooling2D (name='avg_flatten') (net)

# output layer
outputs = layers.Dense (units=targets.shape[1], activation='softmax', name='output') (net)

model = models.Model (inputs=inputs, outputs=outputs)
model.summary ()
```

| Layer (type) | Output Shape | Param # |
|------------------------------|----------------------|---------|
| input (InputLayer) | (None, 162, 162, 1) | 0 |
| conv2d_1 (Conv2D) | (None, 162, 162, 8) | 16 |
| conv2d_2 (Conv2D) | (None, 162, 162, 16) | 528 |
| maxp_2 (MaxPooling2D) | (None, 81, 81, 16) | 0 |
| conv2d_3 (Conv2D) | (None, 81, 81, 32) | 2080 |
| maxp_3 (MaxPooling2D) | (None, 40, 40, 32) | 0 |
| conv2d_4 (Conv2D) | (None, 40, 40, 64) | 8256 |
| maxp_4 (MaxPooling2D) | (None, 20, 20, 64) | 0 |
| avg_flatten (GlobalAveragePo | (None, 64) | 0 |
| output (Dense) | (None, 24) | 1560 |
| Total params: 12,440 | | |
| Trainable params: 12,440 | | |
| Non-trainable params: 0 | | |

Model parameter

(metric, loss function)

```
In [13]: # from: Arseny Kravchenko http://arseny.info/2017/f-beta-score-for-keras.html
from keras import backend as K

PARAM_BETA = 1
def fbeta (y_true, y_pred):

    # just in case of hipster activation at the final layer
    y_pred = K.clip (y_pred, 0, 1)

    tp = K.sum (K.round (y_true * y_pred)) + K.epsilon ()
    fp = K.sum (K.round (K.clip (y_pred - y_true, 0, 1)))
    fn = K.sum (K.round (K.clip (y_true - y_pred, 0, 1)))

    precision = tp / (tp + fp)
    recall = tp / (tp + fn)

    beta_squared = PARAM_BETA ** 2
    return (beta_squared + 1) * (precision * recall) / (beta_squared * precision + recall)
```

```
In [14]: from keras import optimizers, losses

PARAM_LR = 0.0001
opt_sgd = optimizers.SGD (lr=PARAM_LR, momentum=0.8)
opt_adamax = optimizers.Adamax (lr=PARAM_LR, beta_1=0.9, beta_2=0.999, epsilon=None, decay=0.0)

model.compile (optimizer=opt_sgd, loss=losses.mean_squared_error, metrics=[fbeta])
```

Model Training and Evaluation

Model training

```
In [15]: import os
         from keras import callbacks

PARAM_MAX_EPOCHS = 100 # PARAM: number of model-fit runs
PARAM_N_BATCH = 10 # PARAM: number of input samples for one feedfwd-backprop step

checkpointer = callbacks.ModelCheckpoint (
    filepath=os.path.join ('model','model.w.best.h5'),
    verbose=1,
    save_best_only=True)

history = model.fit (X_train, y_train,
                    epochs=PARAM_MAX_EPOCHS, batch_size=PARAM_N_BATCH, validation_split=0.1, shuffle=True,
                    callbacks=[checkpointer], verbose=1)
```

```
Train on 194 samples, validate on 22 samples
Epoch 1/100
194/194 [=====] - 7s 35ms/step - loss: 0.0516 - fbeta: 0.0331 - val_loss: 0.0516
- val_fbeta: 2.3849e-08

Epoch 00001: val_loss improved from inf to 0.05162, saving model to model/model.w.best.h5
Epoch 2/100
194/194 [=====] - 6s 32ms/step - loss: 0.0492 - fbeta: 0.0155 - val_loss: 0.0495
- val_fbeta: 2.6446e-08

Epoch 00002: val_loss improved from 0.05162 to 0.04949, saving model to model/model.w.best.h5
Epoch 3/100
194/194 [=====] - 6s 32ms/step - loss: 0.0478 - fbeta: 1.9171e-08 - val_loss: 0.
0484 - val_fbeta: 2.7273e-08

Epoch 00003: val_loss improved from 0.04949 to 0.04838, saving model to model/model.w.best.h5
Epoch 4/100
194/194 [=====] - 6s 32ms/step - loss: 0.0470 - fbeta: 2.0244e-08 - val_loss: 0.
0477 - val_fbeta: 2.7273e-08

Epoch 00004: val_loss improved from 0.04838 to 0.04773, saving model to model/model.w.best.h5
Epoch 5/100
194/194 [=====] - 6s 33ms/step - loss: 0.0466 - fbeta: 2.0619e-08 - val_loss: 0.
0473 - val_fbeta: 2.7273e-08

Epoch 00005: val_loss improved from 0.04773 to 0.04730, saving model to model/model.w.best.h5
Epoch 6/100
194/194 [=====] - 6s 33ms/step - loss: 0.0463 - fbeta: 2.0619e-08 - val_loss: 0.
0470 - val_fbeta: 2.7273e-08

Epoch 00006: val_loss improved from 0.04730 to 0.04701, saving model to model/model.w.best.h5
Epoch 7/100
194/194 [=====] - 6s 33ms/step - loss: 0.0460 - fbeta: 2.0619e-08 - val_loss: 0.
0468 - val_fbeta: 2.7273e-08

Epoch 00007: val_loss improved from 0.04701 to 0.04678, saving model to model/model.w.best.h5
Epoch 8/100
194/194 [=====] - 6s 33ms/step - loss: 0.0458 - fbeta: 2.0619e-08 - val_loss: 0.
0466 - val_fbeta: 2.7273e-08

Epoch 00008: val_loss improved from 0.04678 to 0.04658, saving model to model/model.w.best.h5
Epoch 9/100
194/194 [=====] - 6s 32ms/step - loss: 0.0456 - fbeta: 2.0619e-08 - val_loss: 0.
0464 - val_fbeta: 2.7273e-08

Epoch 00009: val_loss improved from 0.04658 to 0.04643, saving model to model/model.w.best.h5
Epoch 10/100
194/194 [=====] - 6s 32ms/step - loss: 0.0455 - fbeta: 2.0619e-08 - val_loss: 0.
0463 - val_fbeta: 2.7273e-08

Epoch 00010: val_loss improved from 0.04643 to 0.04628, saving model to model/model.w.best.h5
Epoch 11/100
194/194 [=====] - 6s 32ms/step - loss: 0.0453 - fbeta: 2.0619e-08 - val_loss: 0.
0461 - val_fbeta: 2.7273e-08

Epoch 00011: val_loss improved from 0.04628 to 0.04615, saving model to model/model.w.best.h5
Epoch 12/100
194/194 [=====] - 6s 33ms/step - loss: 0.0452 - fbeta: 2.0619e-08 - val_loss: 0.
0460 - val_fbeta: 2.7273e-08

Epoch 00012: val_loss improved from 0.04615 to 0.04603, saving model to model/model.w.best.h5
Epoch 13/100
194/194 [=====] - 6s 33ms/step - loss: 0.0451 - fbeta: 2.0619e-08 - val_loss: 0.
0459 - val_fbeta: 2.7273e-08

Epoch 00013: val_loss improved from 0.04603 to 0.04592, saving model to model/model.w.best.h5
Epoch 14/100
194/194 [=====] - 6s 33ms/step - loss: 0.0450 - fbeta: 2.0619e-08 - val_loss: 0.
0458 - val_fbeta: 2.7273e-08

Epoch 00014: val_loss improved from 0.04592 to 0.04580, saving model to model/model.w.best.h5
Epoch 15/100
194/194 [=====] - 6s 33ms/step - loss: 0.0449 - fbeta: 2.0619e-08 - val_loss: 0.
0457 - val_fbeta: 2.7273e-08

Epoch 00015: val_loss improved from 0.04580 to 0.04570, saving model to model/model.w.best.h5
Epoch 16/100
194/194 [=====] - 6s 33ms/step - loss: 0.0448 - fbeta: 2.0619e-08 - val_loss: 0.
0456 - val_fbeta: 2.7273e-08
```


Model evaluation and comparison

```
In [16]: print (history.history.keys())

dict_keys(['fbeta', 'val_loss', 'val_fbeta', 'loss'])
```

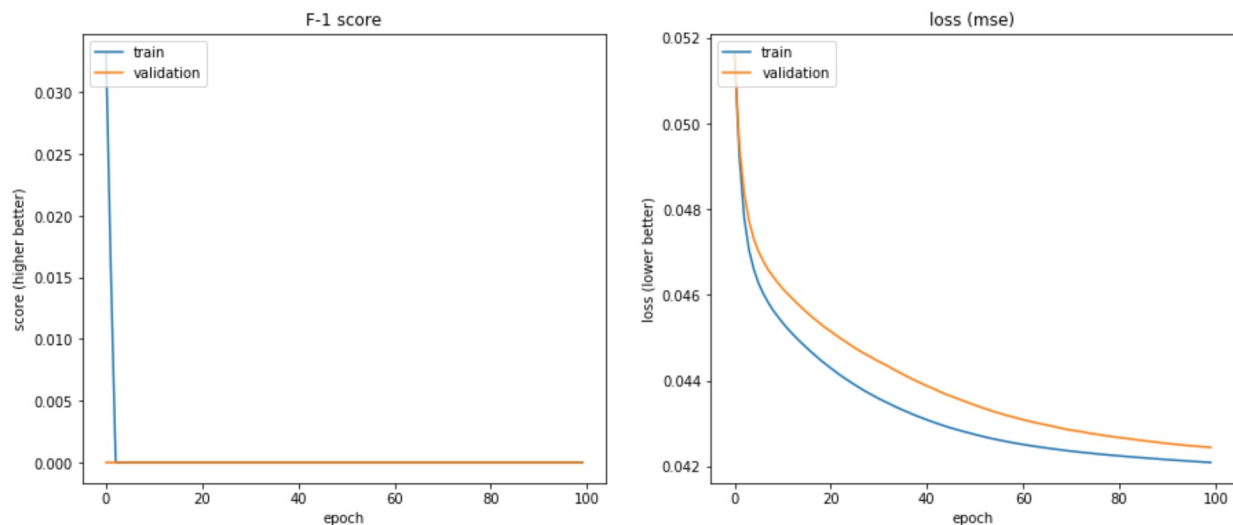
```
In [26]: import matplotlib.pyplot as plt
%matplotlib inline

fig, axs = plt.subplots (1, 2)

# summarize history for accuracy
axs[0].plot (history.history['fbeta'])
if 'val_fbeta' in history.history:
    axs[0].plot (history.history['val_fbeta'])
axs[0].set (xlabel='epoch', ylabel='score (higher better)', title='F-{} score'.format (PARAM_BETA))
axs[0].legend (['train', 'validation'], loc='upper left')

# summarize history for loss
axs[1].plot (history.history['loss'])
if 'val_loss' in history.history:
    axs[1].plot (history.history['val_loss'])
axs[1].set (xlabel='epoch', ylabel='loss (lower better)', title='loss (mse)')
axs[1].legend (['train', 'validation'], loc='upper left')

fig.set_size_inches ((15., 6.), forward=True)
plt.show ()
```



discussion

Above graphs show the F-beta score per epoch with $\beta = 1$ on the left and the *loss per epoch*, calculated by the mean squared error (mse) on the right.

loss per epoch:

- gradient steps start with a loss of 0.052, end by 0.042 and show a smooth concave curve. The curve couldn't be better except a faster drop in the first 10 epochs.
- the worse: mse after 1st epoch = 0.052 - the CNN learns very slow and in tiny steps (1st/2nd epoch: $0.052 - 0.049 = 0.003$)

F-beta score per epoch

- evaluation metric immediately drops to zero after some epochs - the CNN doesn't learn anything yet

reasons / todo*input data*

(1) The used dataset only has 240 samples for training, validation and test. This is by far nothing for the CNN.

Todo: retrieve more samples for the dataset

(2) A quick look at random spectrograms show kind of chaotic information - as a human being it is hard to tell if there's any structure behind each key-mode pair. This may apply to the CNN too.

Todo: find additional filter techniques / methods to clearly bring out structures for the CNN

(3) Songs can change in key over their whole length.

Todo: take appropriate sample of a song - omit bridges, refrains, silent passages, noisy songs

—

model training

The model was trained for 100 epochs, each in batches of 10 samples per feedfwd-backprop step. To make sure that the architecture is well suited, more epochs shall be run.

Todo: increase epochs, change batch size

—

model architecture

Todo: To better understand the insight of the CNN, visualize the filter of the convolutions. May there be enlightenment what kind of architecture works best.

compare learning algorithm to benchmarks

[i] below statements can be run without executing the whole notebook

Therefor, go to and execute [load learning algorithm](#)

TODO TODO TODO

```
In [34]: score = model.evaluate (X_test, y_test, verbose=1)
         print (score)

24/24 [=====] - 0s 16ms/step
[0.042050689458847046, 8.333333134658005e-09]
```

Misc**save learning algorithm**

```
In [30]: # serialization of model architecture
         import os

         save_name = os.path.join ('model', 'model.arch.yaml')

         print ('>>> saving model...', end=' ', flush=True)
         yaml_string = model.to_yaml ()
         with open (save_name, 'w') as yaml_file:
             yaml_file.write (yaml_string)
         print ('done')

>>> saving model... done
```

load learning algorithm

```
In [33]: # load model architecture
from keras import models

print ('>>> loading and compiling model...', end=' ', flush=True)
with open (save_name, 'r') as yaml_file:
    yaml_string = yaml_file.read ()
model = models.model_from_yaml (yaml_string)
model.compile (optimizer=opt_sgd, loss=losses.mean_squared_error, metrics=[fbeta])
print ('done')

# load best weights
print ('>>> loading best weights into model...', end=' ', flush=True)
model.load_weights (os.path.join ('model', 'model.w.best.h5'))
print ('done')

>>> loading and compiling model... done
>>> loading best weights into model... done
```

```
In [20]: idx = 20
test_file = src_spectro_data['filenames'][idx]

test_spectro = path_to_tensor (test_file)
test_pred = model.predict (test_spectro)

print (test_file)
print ('y_true', src_spectro_data['target_names'][src_spectro_data['target'][idx]])
print ('y_pred', src_spectro_data['target_names'][test_pred.argmax ()])

src_spectro/4-0/TROEPPK128F92F33EC.png
y_true 4-0
y_pred 6-1
```

Obsolete

drawbacks (known, unresolvable issues)

(WRONG) *music keys vs CNN key classes*

See [Chuan, Ching-Hua & Chew, Elaine. \(2018\). Audio onset detection using machine learning techniques: the effect and applicability of key and tempo information. \(https://www.researchgate.net/publication/228963946_Audio_onset_detection_using_machine_learning_techniques_the_effect_and_applicability_of_key_and_tempo_information\)](https://www.researchgate.net/publication/228963946_Audio_onset_detection_using_machine_learning_techniques_the_effect_and_applicability_of_key_and_tempo_information), p. 18

The spectrograms show a pitch range given by the [Scientific Pitch Notation \(https://en.wikipedia.org/wiki/Scientific_pitch_notation#Table_of_note_frequencies\)](https://en.wikipedia.org/wiki/Scientific_pitch_notation#Table_of_note_frequencies). By that the range of notes goes from $C_{-1} = 0_{MIDI}$ up to $G_9 = 127_{MIDI}$.

Each note can be the tonic of a music key - for example the key 'C major' exists 11 times (octave -1 to 9). Thus the information of 128 keys is now squeezed into 24 key classes.

...